

CS411 Database Systems

Spring 2015, Prof. Chang

Department of Computer Science
University of Illinois at Urbana-Champaign

Final Examination
May 8, 2015
Time Limit: 180 minutes

- Print your name and NetID below. In addition, print your NetID in the upper right corner of every page.

Name: _____ **NetID:** _____

- Including this cover page, this exam booklet contains **17** pages. Check if you have missing pages.
- The exam is closed book and closed notes. You are allowed to use non-programmable calculators. No other electronic devices are permitted. Any form of cheating on the examination will result in a zero grade.
- Please write your solutions in the spaces provided on the exam. You may use the blank areas and backs of the exam pages for scratch work.
- Please make your answers clear and succinct; you will lose credit for verbose, convoluted, or confusing answers. *Simplicity does count!*
- Each problem has different weight, as listed below– So, plan your time accordingly.

| Problem | 1 | 2 | 3 | 4 | 5 | 6 | Total |
|---------|----|----|----|----|----|----|-------|
| Points | 54 | 26 | 12 | 15 | 18 | 25 | 150 |
| Score | | | | | | | |

Problem 1 (54 points) Misc. Concepts

For each of the following statements:

- for true/false choices, indicate whether it is *TRUE* or *FALSE* by **circling** your choice, and provide an **explanation** to justify;
- for short answer questions, provide a brief **answer** showing your work.

You will get 3 points for each correct answer with correct explanations, and **no penalty (of negative points) for wrong answers**. However, for a question requiring explanation, if you give us incorrect/missing explanation, you will get 0 point even when the answer is correct!

Note: Questions 1-3 use the following two tables $R(A, B, C)$ and $S(C, D)$.

Table R

| A | B | C |
|-------|---|---|
| ===== | | |
| 1 | 2 | 8 |
| 1 | 4 | 7 |
| 2 | 5 | 5 |
| 2 | 4 | 3 |

Table S

| C | D |
|-------|---------|
| ===== | |
| 2 | (1, 2) |
| 3 | (5, 6) |
| 5 | (6, 9) |
| 7 | (5, 11) |

(1) Answer: True False

$\{A, B\}$ can be a key for table R .

\Rightarrow Explain: _____

(2) Write down the output of $(\pi_A R) \bowtie_{R.A=S.C} S$.

\Rightarrow Answer: _____

(3) For a natural join between the tables R and S , i.e., $R \bowtie S$, what is the number of rows and columns of the resulting table?

\Rightarrow Answer: _____

(4) Answer: True False

In a linear hash table, the number of buckets n must be a power of 2.

\Rightarrow Explain:_____

(5) Assume join is a commutative binary operator, i.e., $A \bowtie B$ is equivalent to $B \bowtie A$. How many distinct query trees are there for query $A \bowtie B \bowtie C$?

\Rightarrow Answer:_____

(6) Answer: True False

A left-deep query tree can support not only materialization but also pipelining in query processing.

\Rightarrow Explain:_____

(7) Answer: True False

Rule-based query optimization does not always give us the optimal query plan; in contrast, cost-based query optimization always gives us the optimal query plan.

\Rightarrow Explain:_____

(8) Answer: True False

If we execute all possible query plans, we are able to make a correct query plan decision because we know the actual cost for each plan. So, this approach is more practical than rule-based query optimization.

\Rightarrow Explain:_____

(9) Answer: True False

A REDO logging system must write a log entry after each update of a database value on disk.

\Rightarrow Explain:_____

(10) Among the four “ACID” properties achieved by transaction management, which two properties are guaranteed by failure recovery?

\Rightarrow Answer:_____

(11) Answer: True False

To perform a non-quiescent checkpoint, the database needs to stop accepting new connections.

\Rightarrow Explain:_____

(12) Answer: True False

In case of UNDO logging, $\langle T1, A, 1 \rangle$ means that transaction T1 has changed the database element A, and its new value is 1.

\Rightarrow Explain:_____

(13) Answer: True False

A relation with two attributes, e.g., $R(A, B)$, with an unknown set of functional dependencies, must necessarily be in 3NF.

\Rightarrow Explain:_____

(14) Answer: True False

Functional dependencies can be inferred by observing data.

\Rightarrow Explain:_____

(15) Answer: True False

In ER modeling, every entity must own a minimal set of uniquely identifying attributes, which is called the entity's primary key.

\Rightarrow Explain:_____

(16) Consider two relations about people who have lived in Champaign and Urbana:

ChampaignLog(SSN, name, occupation, info), *UrbanaLog*(SSN, name, occupation, info).

Apply algebraic laws on the following relational algebra expression to optimize the cost. Make sure selection and projection operations are executed as soon and in as small scope as possible.

$\pi_{SSN, occupation, name}(\sigma_{occupation='professor'}(ChampaignLog - UrbanaLog))$

\Rightarrow Answer:_____

- (17) Convert the E/R diagram below to a relational database schema, using one of the following approaches: E/R, OO and Nulls. You only need to choose one of the methods, but you must indicate which method you use. Failure to do so will lead to a zero score for this sub question.

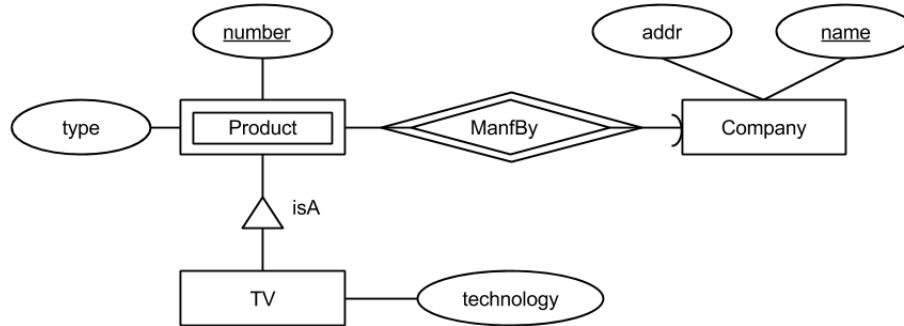


Figure 1: E/R diagram

⇒ Answer: _____

- (18) Consider two relations of UIUC student/enrollment records: *Student*(netid, department, name, state) and *Enrollment*(netid, courseid, grade). There are 40,000 tuples in *Student* from 100 departments and 50 states, which include every student at UIUC. There are 200,000 tuples in *Enrollment*. Estimate the size, in terms of number of tuples, for

$$(\sigma_{\text{department} = \text{"CS"}} \textit{Student}) \bowtie \textit{Enrollment}.$$

⇒ Answer: _____

Problem 2 (*22 points*) Query Languages

- (1) Answer the following questions based on the database schema provided below for a class registration system.

Class(ClassID, ClassName, InstructorID, Time, Location)

Instructor(InstructorID, Name)

Student(StudentID, Name, Major, GPA, TotalHours)

Register(StudentID, ClassID, PointGrade, CreditHours)

- (a) Write a query, in *relational algebra*, to return the names of students with highest GPA. There may be more than one student with the highest GPA; in such case, return all such students. (*4 points*)

- (b) Write a query, in *SQL*, to return the names of students who have registered for class taught by 'Angrave, L.'. (*4 points*)

Class(ClassID, ClassName, InstructorID, Time, Location)

Instructor(InstructorID, Name)

Student(StudentID, Name, Major, GPA, TotalHours)

Register(StudentID, ClassID, PointGrade, CreditHours)

(c) Write a query, in *SQL*, to return the names of instructors who teach more than 3 classes. The output should be in descending order by number of classes. (4 points)

(d) In *SQL*, define a foreign key constraint to ensure whenever a student record is deleted from the *Student* table, his/her records are deleted from the *Register* table as well. (5 points)

Class(ClassID, ClassName, InstructorID, Time, Location)

Instructor(InstructorID, Name)

Student(StudentID, Name, Major, GPA, TotalHours)

Register(StudentID, ClassID, PointGrade, CreditHours)

(e) Write a trigger, in *SQL*, that updates GPA for a student whenever there is an update on *Register* table. GPA is calculated using formula below: (5 points)

$$GPA = \sum_{\text{class} \in \text{all classes taken}} \frac{PointGrade(class) * CreditHours(class)}{TotalHours}$$

Problem 3 (16 points) Indexing

- (1) Consider the following B+Tree of order 4 (i.e., $n=4$, each index can hold n keys and $n + 1$ pointers), shown in the figure below:




Figure 2: B+Tree

For the following two questions, **please draw the full tree**. Any incomplete drawing will result in **0** points.

- (a) Show the resulting tree after inserting key 8. (4 points)

- (b) Based on the original tree, show the resulting tree after deleting 13. (4 points)

- (2) Consider indexing the following key values using an extensible hash table. Suppose that we insert the keys in the order of: 1, 15, 21, 35.

The hash function $h(n)$ for key n is $h(n) = n \bmod 16$; i.e., the hash function is the remainder after the key value is divided by 16. Thus, the hash value is a 4-bit value. Assume that each bucket can hold 2 data items.

You have performed this indexing correctly, and have come up with the following table, as shown in figure below:




Figure 3: Extensible Hashing Table

Now insert 18 and 41, in this order, into the hash table into the table above. Redraw the table to reflect the new values. Be sure to indicate the number of bits in the hash value that are used in the array. Also, indicate the “nub” value of each block. Again, you can just draw the final table after inserting all the six keys. (8 points)

(1) If we are using Block-based nested-loop join algorithm for $R \text{ JOIN } S$, what is the proper I/O cost if $B(R) = 8$, $B(S) = 15$, $M = 6$? Feel free to directly apply the corresponding formula. (4 points)

- 12

- (3) If we are using Sort-Merge join algorithm for $R \text{ JOIN } S$, what is the proper I/O cost if $B(R) = 8$, $B(S) = 15$, $M = 6$? Feel free to directly apply the corresponding formula. (4 points)

- (4) If we are using Hash-join algorithm for $R \text{ JOIN } S$, what is the proper I/O cost if $B(R) = 8$, $B(S) = 15$, $M = 6$? Feel free to directly apply the corresponding formula. (3 points)

Problem 5 (18 points) Dynamic Programming

- (1) Assume we have 50 tuples for relation A , i.e., $T(A) = 50$, and 30 tuples for relation B . What is the possible minimum $T(A \text{ JOIN } B)$? Explanation is not required. (1 points)
- (2) For the same question in **problem 5.1**, what is the possible maximum $T(A \text{ JOIN } B)$? Explanation is not required. (1 points)
- (3) Given relations A, B, C, D , assume $T(A) = 20$, $T(B) = 50$, $T(C) = 70$, $T(D) = 10$, and the size estimation heuristic: $\text{size}(R1 \text{ JOIN } R2) = 0.1 * T(R1) * T(R2)$. In order to find out the best query plan for **A JOIN B JOIN C JOIN D**, please fill in the following table. (16 points)

| | Subquery | Size | Cost | Plan |
|----|----------|------|------|------|
| 1 | AB | | | |
| 2 | AC | | | |
| 3 | AD | | | |
| 4 | BC | | | |
| 5 | BD | | | |
| 6 | CD | | | |
| 7 | ABC | | | |
| 8 | ABD | | | |
| 9 | ACD | | | |
| 10 | BCD | | | |
| 11 | ABCD | | | |

Problem 6 (*25 points*) Failure Recovery

Consider the following log sequence.

| Log ID | Log |
|--------|-------------|
| 1 | ⟨START T1⟩ |
| 2 | ⟨T1, A, 1⟩ |
| 3 | ⟨START T2⟩ |
| 4 | ⟨T2, B, 2⟩ |
| 5 | ⟨COMMIT T2⟩ |
| 6 | ⟨T1, B, 2⟩ |
| 7 | ⟨COMMIT T1⟩ |
| 8 | ⟨START T3⟩ |
| 9 | ⟨T3, A, 3⟩ |
| 10 | ⟨START T4⟩ |
| 11 | ⟨T3, B, 4⟩ |
| 12 | ⟨START T5⟩ |
| 13 | ⟨COMMIT T3⟩ |
| 14 | ⟨T4, C, 5⟩ |
| 15 | ⟨COMMIT T4⟩ |
| 16 | ⟨T5, A, 6⟩ |
| 17 | ⟨COMMIT T5⟩ |
| 18 | ⟨START T6⟩ |
| 19 | ⟨T6, A, 8⟩ |
| 20 | ⟨COMMIT T6⟩ |

- 16

Note: For the questions below, assume the given log sequence is a **REDO** log.

- (d) Suppose that we begin nonquiescent checkpointing right after logID 4. In the space below, indicate where *i*) the start checkpointing record would be, and what it would look like; and *ii*) the earliest end checkpoint record would be, and what it would look like. (*4 points*)
- (e) Continue from (d). Suppose the system crashes right after logID 16. If $\langle \text{ENDCKPT} \rangle$ was written to the log, indicate the portion of the log we would need to inspect and which transactions need to be redone. (*4 points*)
- (f) Now, suppose that we begin nonquiescent checkpointing right after logID 4 and the system crashes right after logID 6. If $\langle \text{ENDCKPT} \rangle$ was not written to the log, indicate the portion of the log we would need to inspect and which transactions need to be redone. (*4 points*)